

Machine Learning

Dr. Vittorio Zanella Palacios

vittorio.zanella@upaep.mx

What is Machine Learning?




Arthur Lee Samuel

(December 5, 1901 – July 29, 1990)

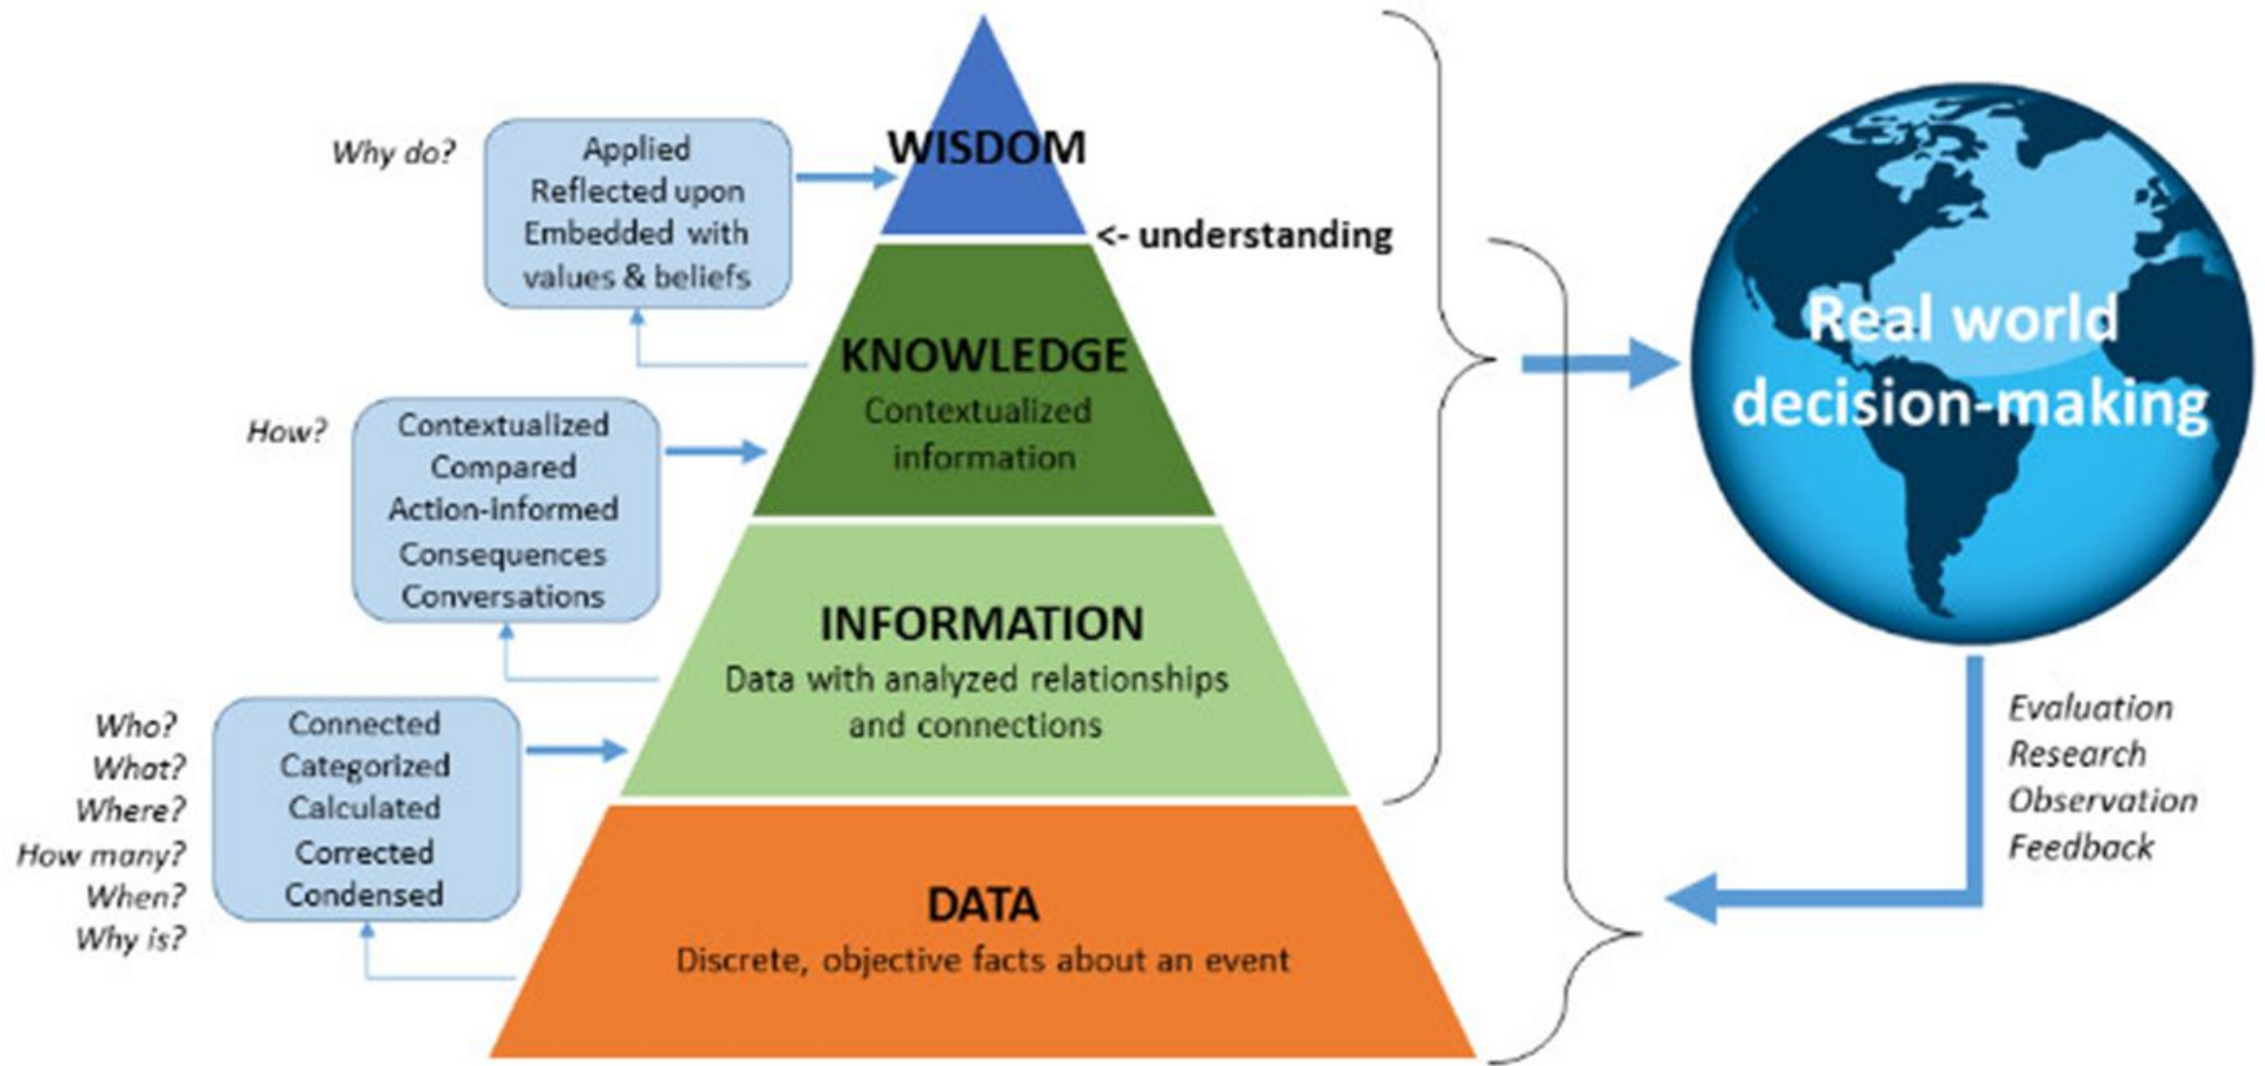
In 1959, *Arthur Samuel* defined machine learning as a “Field of study that gives computers the ability to learn with out being explicitly programmed”

It is a research field at the intersection of statistics, artificial intelligence and computer science






Machine learning is about extracting knowledge from data




2019 *This Is What Happens In An Internet Minute*




This is what happens in a minute on the internet | World Economic Forum



Dilbert.com DilbertCartoonist@gmail.com





2-2-13 ©2013 Scott Adams, Inc./Dist. by Universal Uclick



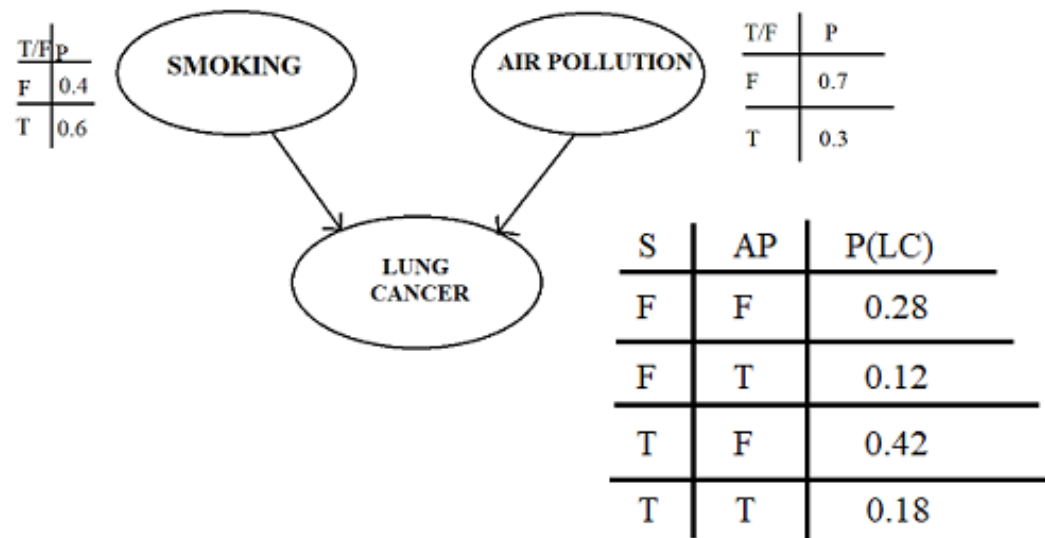

History and relationships to other fields

Already in the early days of AI as an academic discipline, some researchers were interested in having machines learn from data.

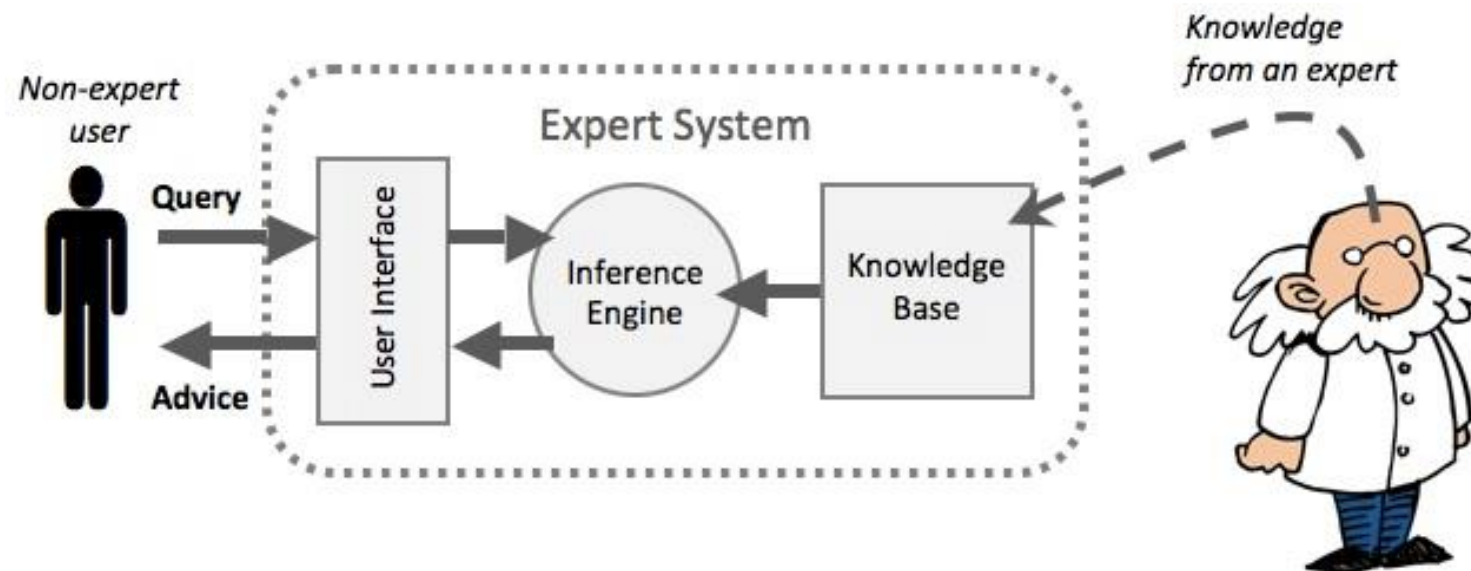


neural networks

Probabilistic reasoning, especially in automated medical diagnosis



By 1980, expert systems had come to dominate AI



But the more statistical line of research was now outside the field of AI proper, in *pattern recognition*.

Neural networks research had been abandoned by AI around the same time.

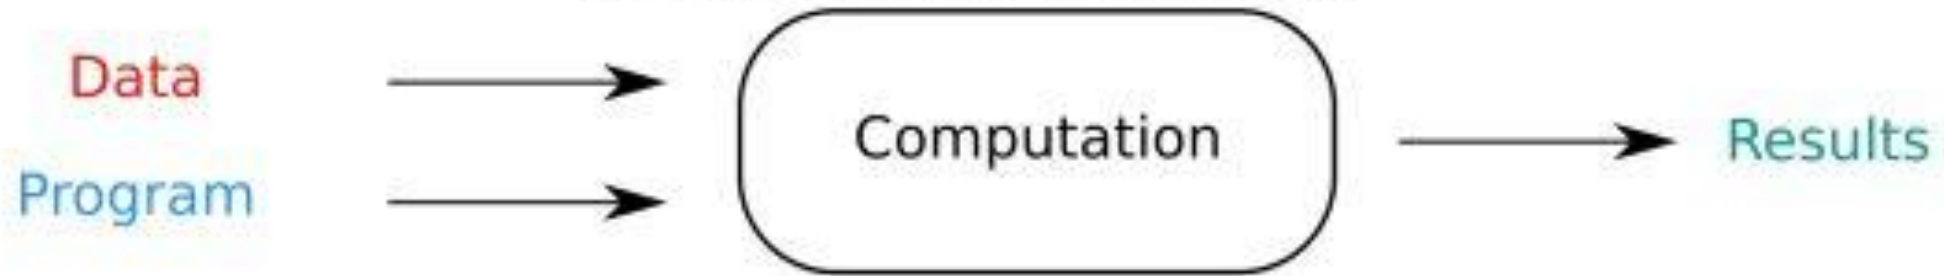
Machine learning, reorganized as a separate field, started to flourish in the 1990s




The field changed its goal from achieving artificial intelligence to tackling solvable problems of a practical nature

It also benefited from the increasing availability of digitized information, and the possibility to distribute that via the internet

Traditional programming



Machine Learning Approach




Why Machine Learning?

Today there is enough computing power

Huge databases exist on virtually every topic

Why Machine Learning Now?



*IDC Digital Universe report, 2014 <http://www.emc.com/infographics/digital-universe-2014.htm>

**Data Scientist: The Sexiest Job of the 21st Century, Oct 2012 <https://hbr.org/2012/10/data-scientist-the-sexiest-job-of-the-21st-century>

To summarize, Machine Learning is great for:

Problems for which existing solutions require a lot of hand-tuning or long lists of rules

Complex problems for which there is no good solution at all using a traditional approach

Fluctuating environments

Getting insights about complex problems and large amounts of data

Problems Machine Learning Can Solve

Homework:

To find and describe two problems that can be solved by machine learning

Types of Machine Learning Systems

There are so many different types of Machine Learning systems that it is useful to classify them in broad categories based on:


- Whether or not they are trained with human supervision (supervised, unsupervised, semisupervised, and Reinforcement Learning)
- Whether or not they can learn incrementally on the fly (online versus batch learning)
- Whether they work by simply comparing new data points to known data points, or instead detect patterns in the training data and build a predictive model, much like scientists do (instance-based versus model-based learning)

Quite possibly the most important part in the machine learning process is understanding the data you are working with and how it relates to the task you want to solve

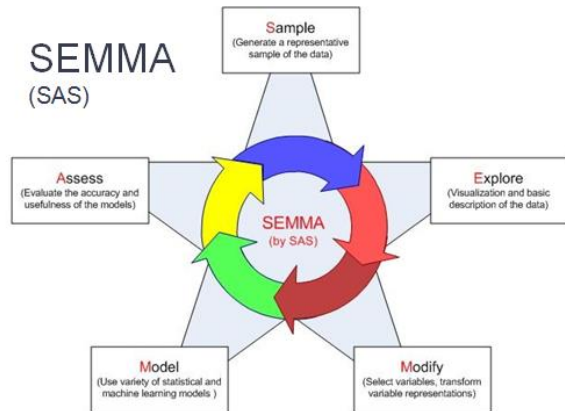
Data Science project management methodologies

From Data to Insight

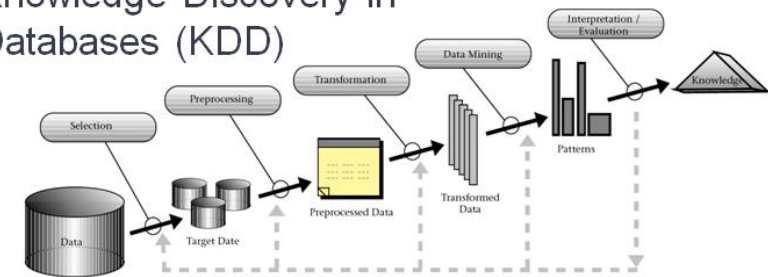
Cross Industry Standard Process for Data Mining (CRISP-DM)
(IBM, Teradata, Daimler AG, NCR Corporation and OHRA)



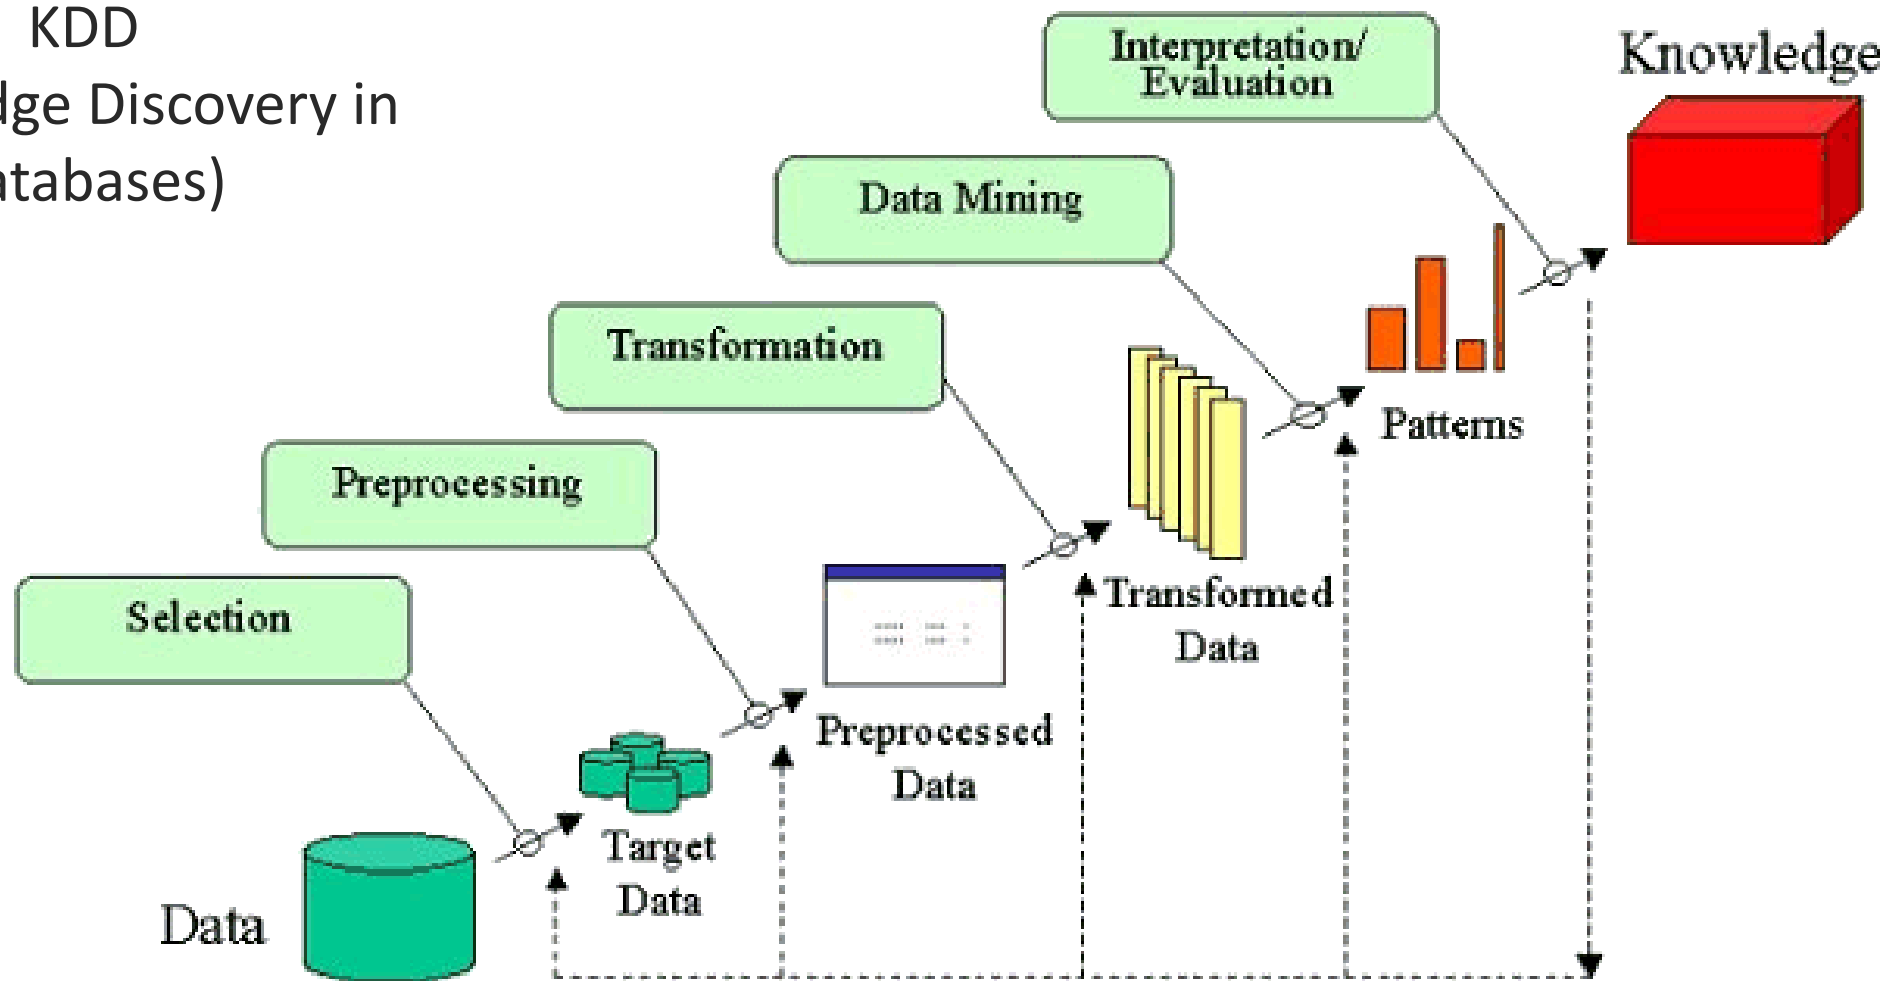
For more information on these methods, see: https://en.wikipedia.org/wiki/Cross_Industry_Standard_Process_for_Data_Mining; <https://en.wikipedia.org/wiki/SEMMA>; https://en.wikipedia.org/wiki/Data_mining




Knowledge Discovery in Databases (KDD)



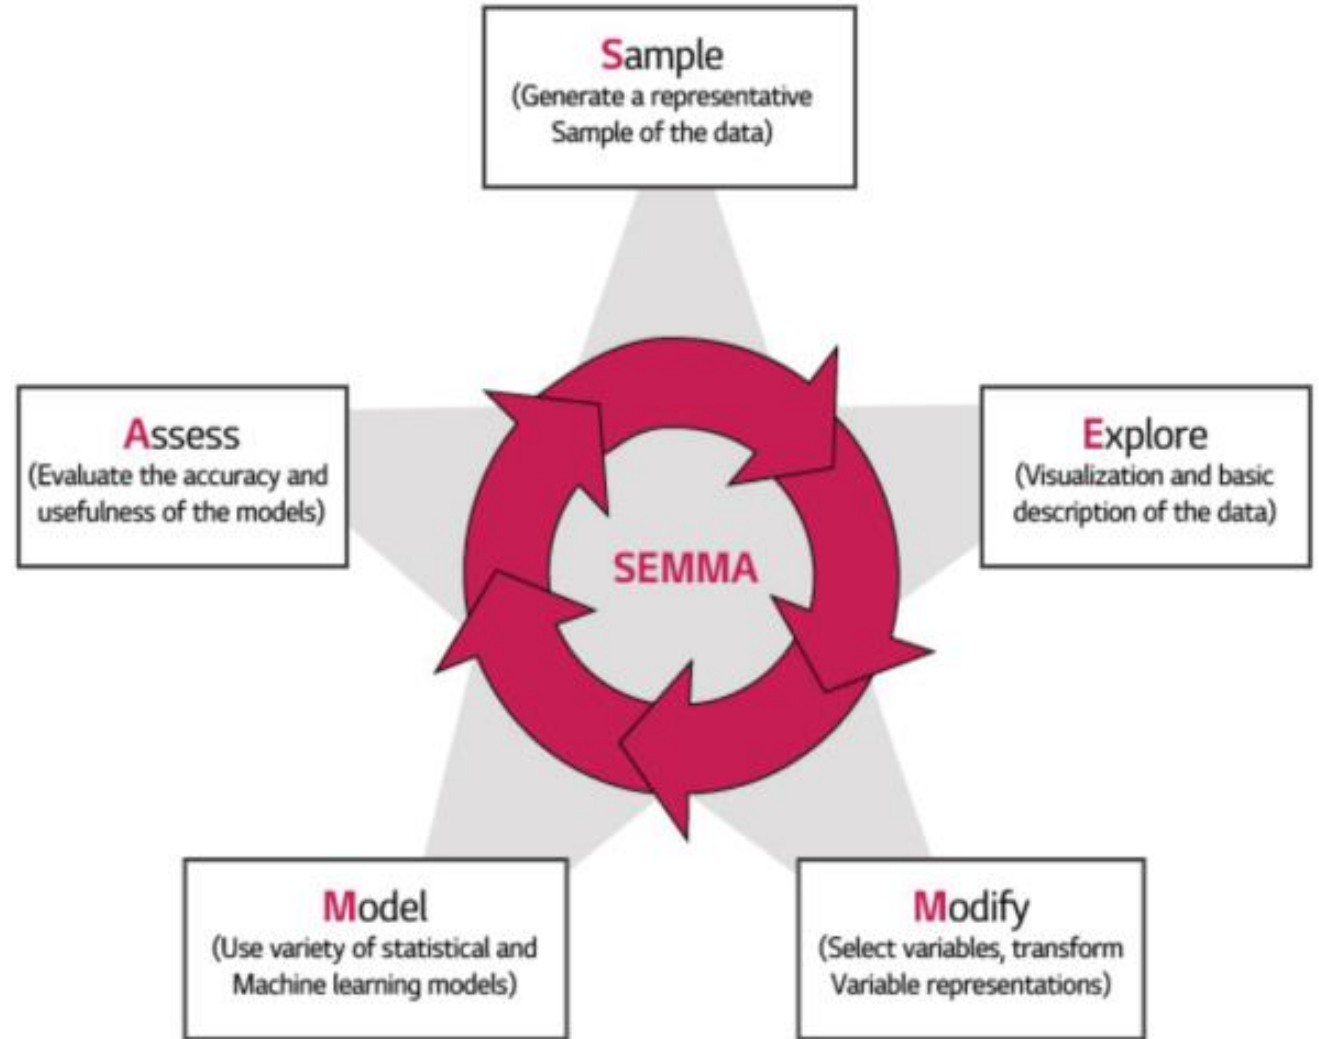
KDD (Knowledge Discovery in Databases)

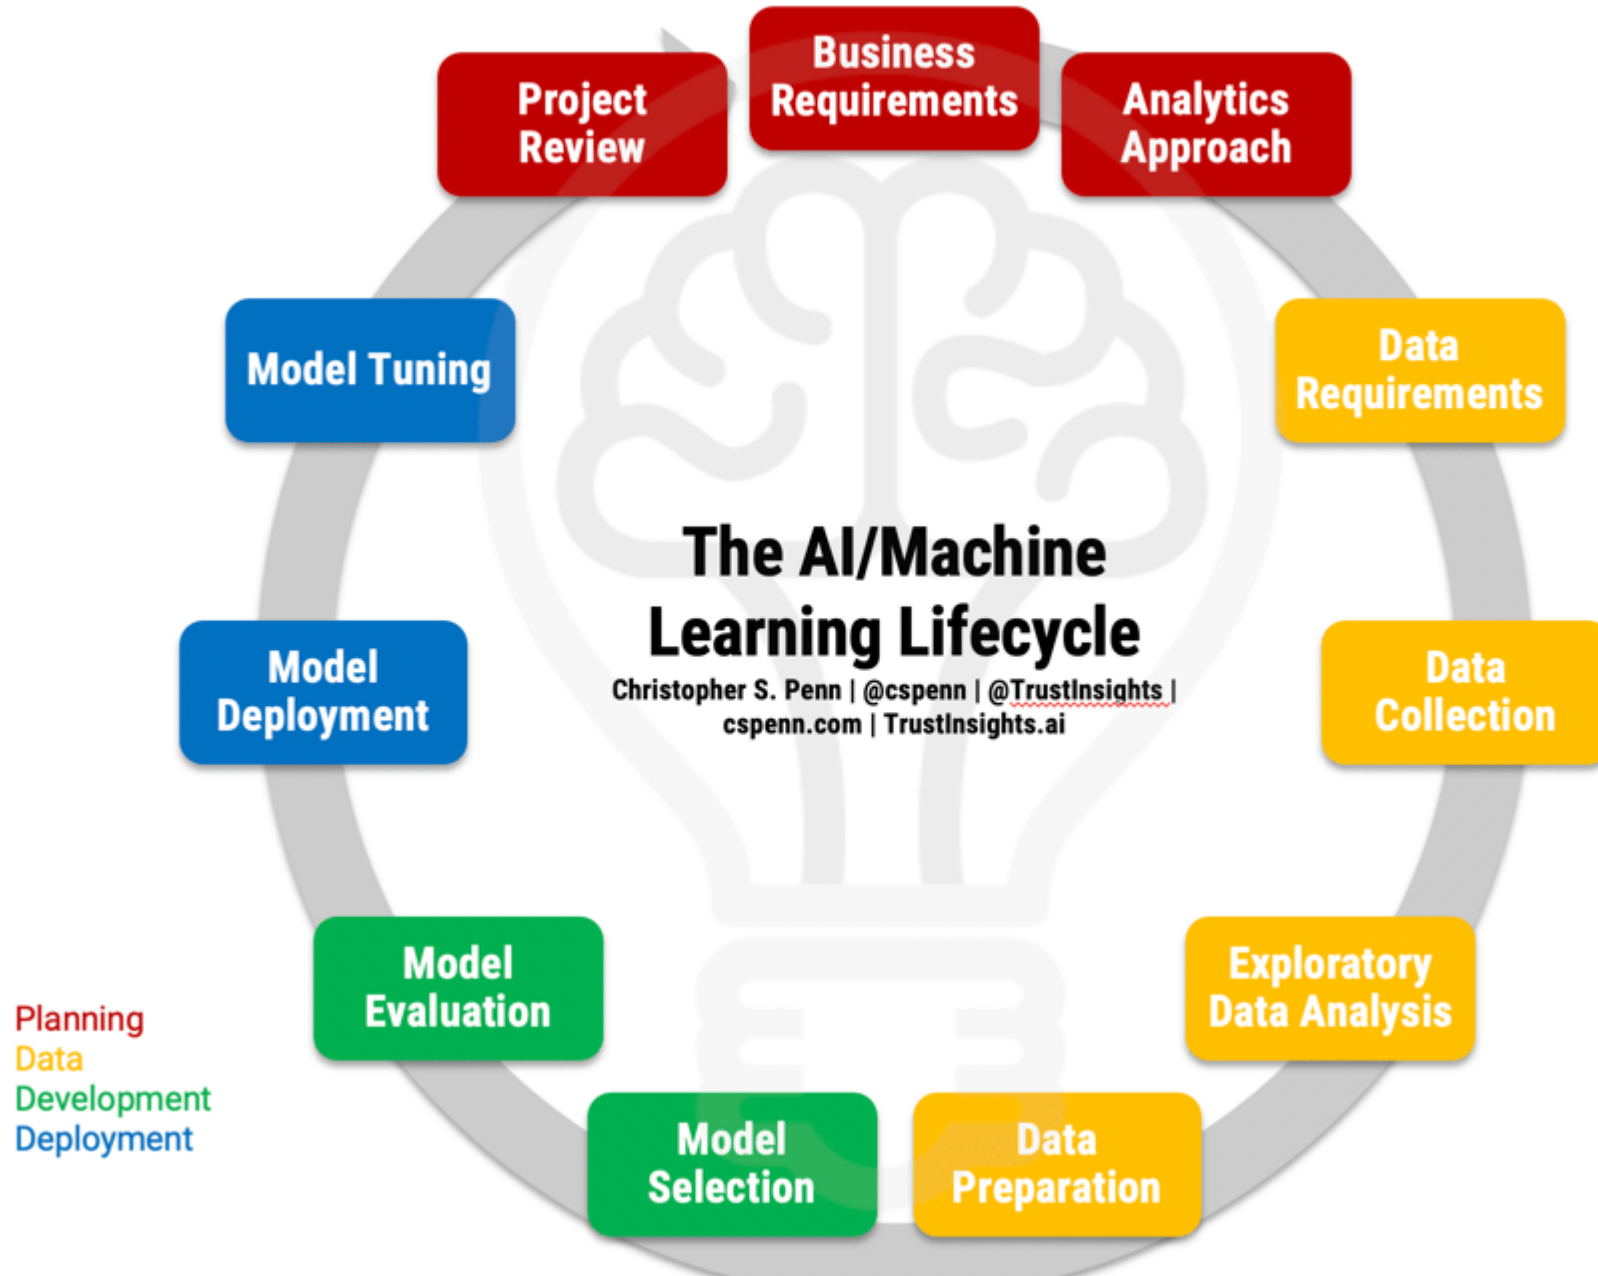


Cross Industry Standard Process for Data Mining (CRISP-DM)



SEMMA (Sample, Explore, Modify, Model and Access)











While you are building a machine learning solution, you should answer, or at least keep in mind, the following questions:

- What question(s) am I trying to answer?
- Do I think the data collected can answer that question?
- What is the best way to phrase my question(s) as a machine learning problem?
- Have I collected enough data to represent the problem I want to solve?
- What features of the data did I extract, and will these enable the right predictions?
- How will I measure success in my application?
- How will the machine learning solution interact with other parts of my research or business product?


Linear regression
Logistic Regression



K-nearest neighbor
Decision tree
Neural Nets



K-Means



Principal Analysis Components (PCA)